

Diabetes Prediction

Importing libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from math import sqrt
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
from sklearn.metrics import accuracy_score
import warnings
from tabulate import tabulate
warnings.simplefilter('ignore')
```

Data Preprocessing

```
df = pd.read_csv('diabetes-vid.csv')
df.head(5)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	dead
1	0.351	31	alive
2	0.672	32	dead
3	0.167	21	alive
4	2.288	33	dead

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 768 entries, 0 to 767
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Pregnancies	768 non-null	int64
1	Glucose	768 non-null	int64
2	BloodPressure	768 non-null	int64
3	SkinThickness	768 non-null	int64
4	Insulin	768 non-null	int64
5	BMI	768 non-null	float64
6	DiabetesPedigreeFunction	768 non-null	float64
7	Age	768 non-null	int64
8	Outcome	768 non-null	object

```
dtypes: float64(2), int64(6), object(1)
```

```
memory usage: 54.1+ KB
```

```
df.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness
count	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458
std	3.369578	31.972618	19.355807	15.952218
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000
75%	6.000000	140.250000	80.000000	32.000000
max	17.000000	199.000000	122.000000	99.000000

	BMI	DiabetesPedigreeFunction	Age
count	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885
std	7.884160	0.331329	11.760232
min	0.000000	0.078000	21.000000
25%	27.300000	0.243750	24.000000
50%	32.000000	0.372500	29.000000
75%	36.600000	0.626250	41.000000
max	67.100000	2.420000	81.000000

```
df.isnull().sum()
```

```

Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction  0
Age              0
Outcome          0
dtype: int64

df.skew()

Pregnancies      0.901674
Glucose           0.173754
BloodPressure    -1.843608
SkinThickness     0.109372
Insulin           2.272251
BMI              -0.428982
DiabetesPedigreeFunction  1.919911
Age              1.129597
dtype: float64

df.mean()

Pregnancies      3.845052
Glucose          120.894531
BloodPressure     69.105469
SkinThickness     20.536458
Insulin           79.799479
BMI              31.992578
DiabetesPedigreeFunction  0.471876
Age              33.240885
dtype: float64

df['Outcome'] = df['Outcome'].apply(lambda x: 1 if x != 'alive' else 0)

```

Exploratory Data Analysis

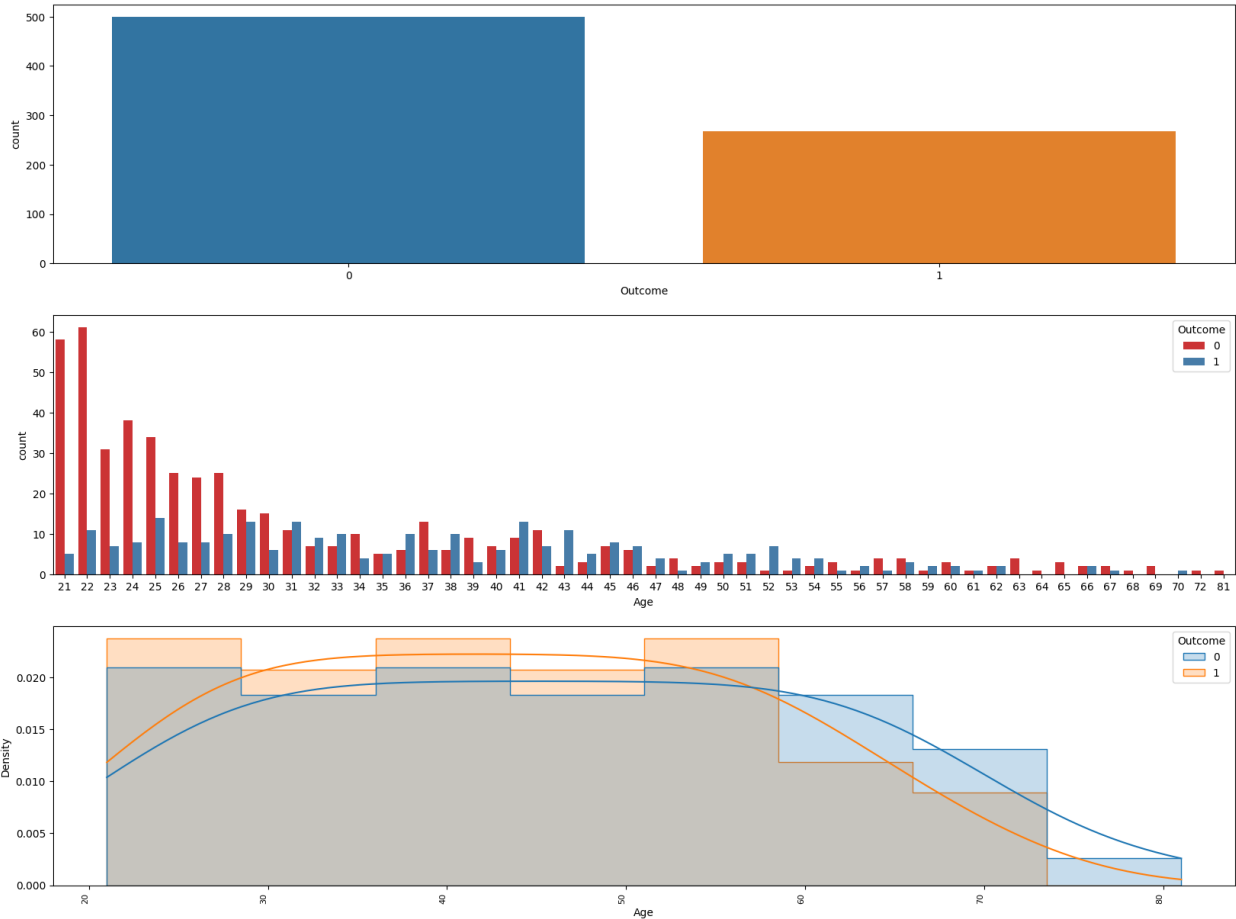
```

fig, ax = plt.subplots(3, 1, figsize=(20, 15))

outcome_count = df.groupby(['Age', 'Outcome'])['Outcome'].count()

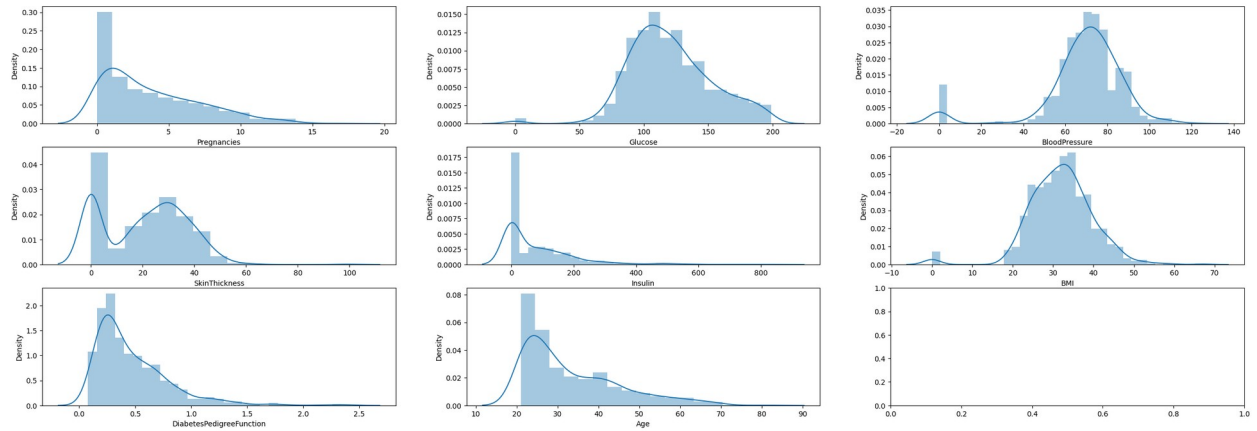
sns.countplot(x=df['Outcome'], ax=ax[0])
sns.countplot(data=df, x='Age', hue='Outcome',
palette='Set1', ax=ax[1])
sns.histplot(data=outcome_count, x='Age', hue='Outcome',
element='step', stat='density', kde = True ,common_norm=False, ax=ax[2])
plt.xticks(rotation=90, ha='right', fontsize=8)
plt.show()

```




```
fig, ax = plt.subplots(3, 3, figsize=(30, 10))

sns.distplot(df['Pregnancies'], ax=ax[0,0])
sns.distplot(df['Glucose'], ax=ax[0,1])
sns.distplot(df['BloodPressure'], ax=ax[0,2])
sns.distplot(df['SkinThickness'], ax=ax[1,0])
sns.distplot(df['Insulin'], ax=ax[1,1])
sns.distplot(df['BMI'], ax=ax[1,2])
sns.distplot(df['DiabetesPedigreeFunction'], ax=ax[2,0])
sns.distplot(df['Age'], ax=ax[2,1])
plt.show()
```



```
correlation_matrix = df.corr()

plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
fmt='.1f')
plt.title('Matriz de Correlación')
plt.show()
```



```
sns.pairplot(df, hue = 'Outcome')
<seaborn.axisgrid.PairGrid at 0x7e812c727220>
```



Data Scaler

```
array = df.values
X = array[ :, 0:8]
Y = array[ :, 8]
names = df.columns.tolist()

from sklearn.preprocessing import Normalizer

scaler = Normalizer().fit(X)
NormalizedX = scaler.transform(X)

print(names)
print(NormalizedX)
```

```

['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome']
[[0.03355237 0.82762513 0.40262844 ... 0.18789327 0.00350622
0.27960308]
 [0.008424    0.71604034 0.55598426 ... 0.22407851 0.00295683
0.26114412]
 [0.04039768 0.92409698 0.32318146 ... 0.11765825 0.00339341
0.16159073]
 ...
 [0.02691539 0.65135243 0.38758161 ... 0.14103664 0.00131885
0.16149234]
 [0.00665306 0.83828547 0.39918356 ... 0.20025708 0.00232192
0.31269379]
 [0.00791454 0.73605211 0.55401772 ... 0.24060198 0.00249308
0.18203439]]

```

```

normalized_df = pd.DataFrame(NormalizedX, columns=names[:8])
normalized_df['Outcome'] = Y
normalized_df.head(1)

```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin
BMI \					
0	0.033552	0.827625	0.402628	0.195722	0.0
	0.187893				

	DiabetesPedigreeFunction	Age	Outcome
0	0.003506	0.279603	1.0

```

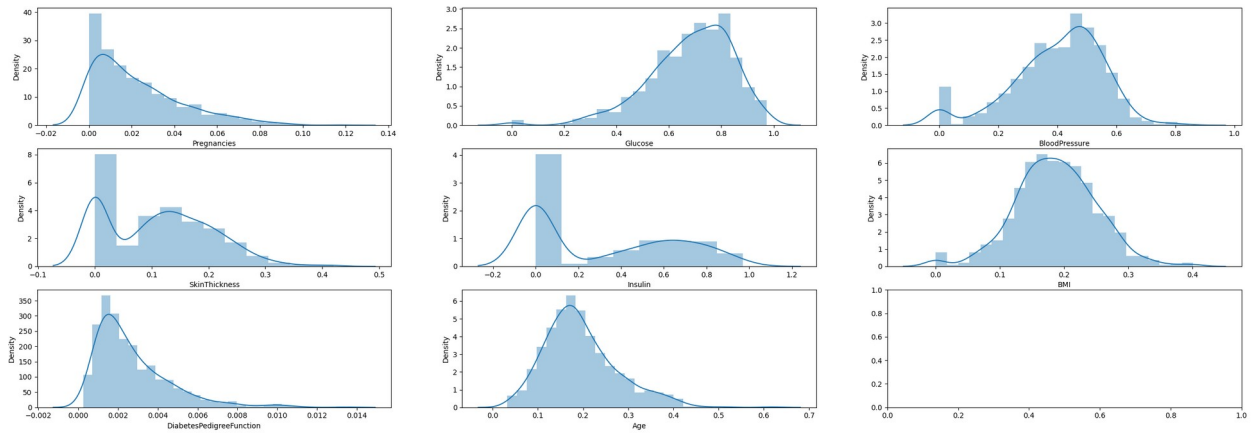
fig, ax = plt.subplots(3, 3, figsize=(30, 10))

```

```

sns.distplot(normalized_df['Pregnancies'],ax=ax[0,0])
sns.distplot(normalized_df['Glucose'],ax=ax[0,1])
sns.distplot(normalized_df['BloodPressure'],ax=ax[0,2])
sns.distplot(normalized_df['SkinThickness'],ax=ax[1,0])
sns.distplot(normalized_df['Insulin'],ax=ax[1,1])
sns.distplot(normalized_df['BMI'],ax=ax[1,2])
sns.distplot(normalized_df['DiabetesPedigreeFunction'],ax=ax[2,0])
sns.distplot(normalized_df['Age'],ax=ax[2,1])
plt.show()

```

```
sns.pairplot(normalized_df, hue = 'Outcome', markers=["o", "D"])
<seaborn.axisgrid.PairGrid at 0x7e8122408df0>
```



Train Test Split

```
from sklearn.model_selection import train_test_split

X_train, X_test, Y_train, Y_test = train_test_split(df.drop('Outcome',
axis = 1), df['Outcome'], test_size=0.25, random_state=42)
```

Modeling and Evaluation

LogisticRegression

```
from sklearn.linear_model import LogisticRegression

model = LogisticRegression()
```

```

model.fit(X_train, Y_train)
model.score(X_train, Y_train)
model_pred = model.predict(X_test)

print(classification_report(Y_test, model_pred))
print("accuracy: ",accuracy_score(Y_test, model_pred))
print("mean_absolute_error: ",mean_absolute_error(Y_test, model_pred))
print("mean_squared_error: ",mean_squared_error(Y_test, model_pred))

```

	precision	recall	f1-score	support
0	0.80	0.77	0.79	123
1	0.62	0.65	0.63	69
accuracy			0.73	192
macro avg	0.71	0.71	0.71	192
weighted avg	0.73	0.73	0.73	192

```

accuracy: 0.7291666666666666
mean_absolute_error: 0.2708333333333333
mean_squared_error: 0.2708333333333333

```

RandomForestClassifier

```

from sklearn.ensemble import RandomForestClassifier

model_RFC =RandomForestClassifier(n_estimators=100, random_state=42)
model_RFC.fit(X_train, Y_train)
model_RFC.score(X_train, Y_train)
model_RFC_Pred = model_RFC.predict(X_test)
print(classification_report(Y_test, model_RFC_Pred))

print("accuracy: ",accuracy_score(Y_test, model_RFC_Pred))
print("mean_absolute_error: ",mean_absolute_error(Y_test,
model_RFC_Pred))
print("mean_squared_error: ",mean_squared_error(Y_test,
model_RFC_Pred))

```

	precision	recall	f1-score	support
0	0.80	0.78	0.79	123
1	0.62	0.65	0.64	69
accuracy			0.73	192
macro avg	0.71	0.72	0.71	192
weighted avg	0.74	0.73	0.74	192

```

accuracy: 0.734375
mean_absolute_error: 0.265625
mean_squared_error: 0.265625

```

SVC

```
from sklearn.svm import SVC

model_SVC = SVC(kernel = 'linear', probability=True ,random_state = 0)
model_SVC.fit(X_train, Y_train)
model_SVC.score(X_train, Y_train)
model_SVC_Pred = model_SVC.predict(X_test)

print(classification_report(Y_test, model_SVC_Pred))
print("accuracy: ",accuracy_score(Y_test, model_SVC_Pred))
print("mean_absolute_error: ",mean_absolute_error(Y_test,
model_SVC_Pred))
print("mean_squared_error: ",mean_squared_error(Y_test,
model_SVC_Pred))
```

	precision	recall	f1-score	support
0	0.79	0.78	0.79	123
1	0.62	0.64	0.63	69
accuracy			0.73	192
macro avg	0.71	0.71	0.71	192
weighted avg	0.73	0.73	0.73	192

```
accuracy: 0.7291666666666666
mean_absolute_error: 0.2708333333333333
mean_squared_error: 0.2708333333333333
```

KNeighborsClassifier

```
from sklearn.neighbors import KNeighborsClassifier

model_NEG = KNeighborsClassifier(n_neighbors=5)
model_NEG.fit(X_train, Y_train)
model_NEG.score(X_train, Y_train)
model_NEG_Pred = model_NEG.predict(X_test)

print(classification_report(Y_test, model_NEG_Pred))
print("accuracy: ",accuracy_score(Y_test, model_NEG_Pred))
print("mean_absolute_error: ",mean_absolute_error(Y_test,
model_NEG_Pred))
print("mean_squared_error: ",mean_squared_error(Y_test,
model_NEG_Pred))
```

	precision	recall	f1-score	support
0	0.74	0.72	0.73	123
1	0.52	0.55	0.54	69

accuracy			0.66	192
macro avg	0.63	0.63	0.63	192
weighted avg	0.66	0.66	0.66	192

```
accuracy: 0.65625
mean_absolute_error: 0.34375
mean_squared_error: 0.34375
```

```
from sklearn.metrics import confusion_matrix
```

```
RLOG = confusion_matrix(Y_test, model_pred)
RMFC = confusion_matrix(Y_test, model_RFC_Pred)
SVC_ = confusion_matrix(Y_test, model_SVC_Pred)
NEG = confusion_matrix(Y_test, model_NEG_Pred)
```


```
fig, axes = plt.subplots(1, 4, figsize=(20, 3))
```

```
sns.heatmap(RLOG, annot=True, cmap='Blues', fmt='g', ax=axes[0])
axes[0].set_title('Confusion Matrix - Logistic Regression')
axes[0].set_xlabel('Predicted')
axes[0].set_ylabel('True')
```

```
sns.heatmap(RMFC, annot=True, cmap='Blues', fmt='g', ax=axes[1])
axes[1].set_title('Confusion Matrix - Random Forest')
axes[1].set_xlabel('Predicted')
axes[1].set_ylabel('True')
```

```
sns.heatmap(SVC_, annot=True, cmap='Blues', fmt='g', ax=axes[2])
axes[2].set_title('Confusion Matrix - SVC')
axes[2].set_xlabel('Predicted')
axes[2].set_ylabel('True')
```

```
sns.heatmap(NEG, annot=True, cmap='Blues', fmt='g', ax=axes[3])
axes[3].set_title('Confusion Matrix - Neighbors')
axes[3].set_xlabel('Predicted')
axes[3].set_ylabel('True')
plt.show()
```



```
fig, ax = plt.subplots(1, 4, figsize=(20, 6))
```

```
sns.distplot(Y_test, label='Real', ax=ax[0], color = 'Green')
sns.distplot(model_pred, label='Predicted', ax=ax[0], color = 'red')
```

```

sns.distplot(Y_test, label='Real', ax=ax[1], color = 'Green')
sns.distplot(model_RFC_Pred, label='Predicted', ax=ax[1], color =
'red')

sns.distplot(Y_test, label='Real', ax=ax[2], color = 'Green')
sns.distplot(model_SVC_Pred, label='Predicted', ax=ax[2], color =
'red')


sns.distplot(Y_test, label='Real', ax=ax[3], color = 'Green')
sns.distplot(model_NEG_Pred, label='Predicted', ax=ax[3], color =
'red')

ax[0].set_title('Logistic Regression')
ax[1].set_title('Random Forest')
ax[2].set_title('SVC')
ax[3].set_title('Neighbors')

ax[0].legend()
ax[1].legend()
ax[2].legend()
ax[3].legend()

<matplotlib.legend.Legend at 0x7d53d16bdf60>

```



```

from sklearn.metrics import roc_curve, auc

probs = model.predict_proba(X_test)[: , 1]
fpr, tpr, thresholds = roc_curve(Y_test, probs)


roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area =
%0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')

```

```
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])

plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc="lower right")
plt.show()
```




```
probs = model_RFC.predict_proba(X_test)[: , 1]
fpr, tpr, thresholds = roc_curve(Y_test, probs)

roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area =
%0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
```



```
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc="lower right")
plt.show()
```



```
probs = model_SVC.predict_proba(X_test)[: , 1]
fpr, tpr, thresholds = roc_curve(Y_test, probs)


roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area =
%0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])

plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
```



```
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc="lower right")
plt.show()
```



```
probs = model_NEG.predict_proba(X_test)[: , 1]
fpr, tpr, thresholds = roc_curve(Y_test, probs)

roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area =
%0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc="lower right")
plt.show()
```

